

Advanced Analytics with Spark

이창언

Ch1. 빅데이터 분석하기

데이터 과학의 어려움

1. 성공적인 분석을 위한 작업의 대부분은 데이터의 전처리 과정에서 이루어진다.
2. 데이터 과학에서 반복(iteration)은 기본적인 과정이다.
3. 잘 돌아가는 모델이 완성되었다고 해서 끝이 아니다.

스파크를 활용해라

스파크는 탐색용 분석 시스템과
운영용 분석 시스템의 간격을 좁힐 수 있다.

Ch2. 스칼라와 스파크를 활용한 데이터 분석

```
import org.apache.spark.sql.{DataFrame, Dataset, Row, SparkSession}
import org.apache.spark.sql.functions._ // for lit(), first(), etc.

// Spark 2.0.2 API for Scala
// sparkSession : https://spark.apache.org/docs/2.0.2/api/scala/index.html#org.apache.spark.sql.SparkSession
// DataSet : https://spark.apache.org/docs/2.0.2/api/scala/index.html#org.apache.spark.sql.Dataset
```

```
import org.apache.spark.sql.{DataFrame, Dataset, Row, SparkSession}
import org.apache.spark.sql.functions._
```

```
case class MatchData(
  id_1: Int,
  id_2: Int,
  cmp_fname_c1: Option[Double],
  cmp_fname_c2: Option[Double],
  cmp_lname_c1: Option[Double],
  cmp_lname_c2: Option[Double],
  cmp_sex: Option[Int],
  cmp_bd: Option[Int],
  cmp_bm: Option[Int],
  cmp_by: Option[Int],
  cmp_plz: Option[Int],
  is_match: Boolean
)
```

```
defined class MatchData
```

Took 1 sec. Last updated by anonymous at January 24 2017, 11:18:43 AM.

```
val spark = SparkSession.builder
  .appName("Intro")
  .getOrCreate
import spark.implicits._ // DataFrame으로 변환을 위한 도움을 준다.
```

```
spark: org.apache.spark.sql.SparkSession = org.apache.spark.sql.SparkSession@2b930bce
import spark.implicits._
```

```
val preview = spark.read.csv("/Users/lee/Documents/sparks/Season2/Advanced_Analytics_With_Spark/linkage/block_1.csv")
preview.show()
```

id_1	id_2	cmp_fname_c1	cmp_fname_c2	cmp_lname_c1	cmp_lname_c2	cmp_sex	cmp_bd	cmp_bm	cmp_by	cmp_plz	is_match
37291	53113	0.8333333333333333	?	1	?	1	1	1	1	0	TRUE
39086	47614	1	?	1	?	1	1	1	1	1	TRUE
70031	70237	1	?	1	?	1	1	1	1	1	TRUE
84795	97439	1	?	1	?	1	1	1	1	1	TRUE
36950	42116	1	?	1	1	1	1	1	1	1	TRUE
42413	48491	1	?	1	?	1	1	1	1	1	TRUE
25965	64753	1	?	1	?	1	1	1	1	1	TRUE
49451	90407	1	?	1	?	1	1	1	1	0	TRUE
39932	40902	1	?	1	?	1	1	1	1	1	TRUE
46626	47940	1	?	1	?	1	1	1	1	1	TRUE
48948	98379	1	?	1	?	1	1	1	1	1	TRUE
4767	4826	1	?	1	?	1	1	1	1	1	TRUE
45463	69659	1	?	1	?	1	1	1	1	1	TRUE
11367	13169	1	?	1	?	1	1	1	1	1	TRUE
10782	89636	1	?	1	?	1	0	1	1	1	TRUE
26206	39147	1	?	1	?	1	1	1	1	1	TRUE
16662	27023	1	1	1	?	1	1	1	1	1	TRUE

Took 2 sec. Last updated by anonymous at January 24 2017 11:18:45 AM

```
preview.schema.foreach(println)
```

```
StructField(_c0,StringType,true)
StructField(_c1,StringType,true)
StructField(_c2,StringType,true)
StructField(_c3,StringType,true)
StructField(_c4,StringType,true)
StructField(_c5,StringType,true)
StructField(_c6,StringType,true)
StructField(_c7,StringType,true)
StructField(_c8,StringType,true)
StructField(_c9,StringType,true)
StructField(_c10,StringType,true)
StructField(_c11,StringType,true)
```

Took 0 sec. Last updated by anonymous at January 24 2017, 11:18:45 AM.

```
val parsed = spark.read
    .option("header", "true") // 파일의 첫 줄을 필드 명으로 사용
    .option("nullValue", "?") // 필드 데이터를 변경( "?" => null )
    .option("inferSchema", "true") // 데이터 타입을 추론한다.
    .csv("/Users/lee/Documents/sparks/Season2/Advanced_Analytics_With_Spark/linkage/block_1.csv")
```

```
parsed: org.apache.spark.sql.DataFrame = [id_1: int, id_2: int ... 10 more fields]
```

Took 2 sec. Last updated by anonymous at January 24 2017, 11:18:47 AM.


```
parsed.show()
```

id_1	id_2	cmp_fname_c1	cmp_fname_c2	cmp_lname_c1	cmp_lname_c2	cmp_sex	cmp_bd	cmp_bm	cmp_by	cmp_plz	is_match
37291	53113	0.8333333333333333	null	1.0	null	1	1	1	1	0	true
39086	47614	1.0	null	1.0	null	1	1	1	1	1	true
70031	70237	1.0	null	1.0	null	1	1	1	1	1	true
84795	97439	1.0	null	1.0	null	1	1	1	1	1	true
36950	42116	1.0	null	1.0	1.0	1	1	1	1	1	true
42413	48491	1.0	null	1.0	null	1	1	1	1	1	true
25965	64753	1.0	null	1.0	null	1	1	1	1	1	true
49451	90407	1.0	null	1.0	null	1	1	1	1	0	true
39932	40902	1.0	null	1.0	null	1	1	1	1	1	true
46626	47940	1.0	null	1.0	null	1	1	1	1	1	true
48948	98379	1.0	null	1.0	null	1	1	1	1	1	true
4767	4826	1.0	null	1.0	null	1	1	1	1	1	true
45463	69659	1.0	null	1.0	null	1	1	1	1	1	true
11367	13169	1.0	null	1.0	null	1	1	1	1	1	true
110782	180636	1.0	null	1.0	null	1	0	1	1	1	true

```
val schema = parsed.schema
schema.foreach(println)
```

FINISHED ▶ ⌂ 📖

```
schema: org.apache.spark.sql.types.StructType = StructType(StructField(id_1,IntegerType,true), StructField(id_2,IntegerType,true), StructField(cmp_fname_c1,DoubleType,true), StructField(cmp_fname_c2,DoubleType,true), StructField(cmp_lname_c1,DoubleType,true), StructField(cmp_lname_c2,DoubleType,true), StructField(cmp_sex,IntegerType,true), StructField(cmp_bd,IntegerType,true), StructField(cmp_bm,IntegerType,true), StructField(cmp_by,IntegerType,true), StructField(cmp_plz,IntegerType,true), StructField(is_match,BooleanType,true))
StructField(id_1,IntegerType,true)
StructField(id_2,IntegerType,true)
StructField(cmp_fname_c1,DoubleType,true)
StructField(cmp_fname_c2,DoubleType,true)
StructField(cmp_lname_c1,DoubleType,true)
StructField(cmp_lname_c2,DoubleType,true)
StructField(cmp_sex,IntegerType,true)
StructField(cmp_bd,IntegerType,true)
StructField(cmp_bm,IntegerType,true)
StructField(cmp_by,IntegerType,true)
StructField(cmp_plz,IntegerType,true)
StructField(is_match,BooleanType,true)
```

```
| parsed.count()
```

```
res57: Long = 574913
```

```
| parsed.cache()
```

```
res59: parsed.type = [id_1: int, id_2: int ... 10 more fields]
```

```
| parsed.groupBy("is_match").count().orderBy($"count".desc).show()
```

```
+-----+-----+
|is_match| count|
+-----+-----+
|  false|572820|
|   true|  2093|
+-----+-----+
```

Took 1 sec. Last updated by anonymous at January 24 2017, 11:18:50 AM.

```
// ex) parsed.describe("id_1","id_2") : 인자값을 넣어 해당 필드만 적용할 수 있다.
parsed.createOrReplaceTempView("parks")
spark.sql("""
    SELECT is_match, COUNT(*) cnt
    FROM parks
    GROUP BY is_match
    ORDER BY cnt DESC
    """).show()
```

is_match	cnt
false	572820
true	2093

Took 2 sec. Last updated by anonymous at January 24 2017, 11:18:51 AM.

```
// describe() : numeric columns, including count, mean, stddev, min, and max의 통계를 리턴해준다.
// ex) parsed.describe("id_1","id_2") : 인자값을 넣어 해당 필드만 적용할 수 있다.
val summary = parsed.describe()
summary.show()
```

FINISHED ▶ ⌂ ⌕

summary	id_1	id_2	cmp_fname_c1	cmp_fname_c2	cmp_lname_c1	cmp_lname_c2	cmp_sex	cmp_bd
cmp_bml	cmp_byl	cmp_plzl						
count	574913	574913	574811	10325	574913	239	574913	574851
mean	33271.962171667714	66564.6636865056	0.7127592938251666	0.8977586763518972	0.31557245780995347	0.32691554145529045	0.9550923357099248	0.22475563232907309
stddev	23622.669425933625	23642.00230967225	0.38892864524635457	0.274257752043053	0.33424946875542494	0.378309202054067	0.20710152240504381	0.41742165872355663
min	1	6	0.0	0.0	0.0	0.0	0	0
max	99894	100000	1.0	1.0	1.0	1.0	1	1

```
summary.select("summary", "cmp_fname_c1", "cmp_fname_c2").show()
```

```
+-----+-----+-----+
|summary|      cmp_fname_c1|      cmp_fname_c2|
+-----+-----+-----+
|  count|          574811|          10325|
|   mean| 0.7127592938251666|0.8977586763518972|
| stddev|0.38892864524635457| 0.274257752043053|
|    min|             0.0|             0.0|
|    max|             1.0|             1.0|
+-----+-----+-----+
```

Took 4 sec. Last updated by anonymous at January 24 2017, 11:18:55 AM.

```
val matches = parsed.where("is_match = true")
val misses = parsed.filter($"is_match" === lit(false))
val matchSummary = matches.describe()
val missSummary = misses.describe()
```

```
matches: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [id_1: int, id_2: int ... 10 more fields]
misses: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [id_1: int, id_2: int ... 10 more fields]
matchSummary: org.apache.spark.sql.DataFrame = [summary: string, id_1: string ... 10 more fields]
missSummary: org.apache.spark.sql.DataFrame = [summary: string, id_1: string ... 10 more fields]
```

Took 5 sec. Last updated by anonymous at January 24 2017, 11:19:00 AM.


```
case class Score(value: Double) {
  def +(oi: Option[Int]) = {
    Score(value + oi.getOrElse(0))
  }
}

def scoreMatchData(md: MatchData): Double = {
  (Score(md.cmp_lname_c1.getOrElse(0.0)) + md.cmp_plz +
    md.cmp_by + md.cmp_bd + md.cmp_bm).value
}

def longForm(desc: DataFrame): DataFrame = {
  import desc.sparkSession.implicitly._
  val schema = desc.schema
  desc.flatMap(row => {
    val metric = row.getString(0)
    (1 until row.size).map(i => (metric, schema(i).name, row.getString(i).toDouble))
  })
  .toDF("metric", "field", "value")
}

def pivotSummary(desc: DataFrame): DataFrame = {
  val lf = longForm(desc)
  lf.groupBy("field").
    pivot("metric", Seq("count", "mean", "stddev", "min", "max")).
    agg(first("value"))
}
```

```
val matchSummary = matches.describe()
val missSummary = misses.describe()
```

```
matchSummary: org.apache.spark.sql.DataFrame = [summary: string, id_1: string ... 10 more fields]
missSummary: org.apache.spark.sql.DataFrame = [summary: string, id_1: string ... 10 more fields]
```

Took 5 sec. Last updated by anonymous at January 24 2017, 11:19:07 AM. (outdated)

```
val matchSummaryT = pivotSummary(matchSummary)
val missSummaryT = pivotSummary(missSummary)
```

```
matchSummaryT: org.apache.spark.sql.DataFrame = [field: string, count: double ... 4 more fields]
missSummaryT: org.apache.spark.sql.DataFrame = [field: string, count: double ... 4 more fields]
```

Took 4 sec. Last updated by anonymous at January 24 2017, 11:19:07 AM. (outdated)

```
matchSummaryT.createOrReplaceTempView("match_desc")
missSummaryT.createOrReplaceTempView("miss_desc")
```

Took 1 sec. Last updated by anonymous at January 24 2017, 11:19:08 AM. (outdated)

```
spark.sql("""
  SELECT a.field, a.count + b.count total, a.mean - b.mean delta
  FROM match_desc a INNER JOIN miss_desc b ON a.field = b.field
  ORDER BY delta DESC, total DESC
""").show()
```

field	total	delta
id_1	1574913.0	1173.1784091823356
cmp_plz	1573618.0	0.9524975516429005
cmp_lname_c2	239.0	0.8136949970410103
cmp_by	1574851.0	0.7763379425859384
cmp_bd	1574851.0	0.7732820129086737
cmp_lname_c1	1574913.0	0.6844795197262346
cmp_bm	1574851.0	0.510834819548174
cmp_fname_c1	1574811.0	0.28531156828536086
cmp_fname_c2	10325.0	0.09900440489032658
cmp_sex	1574913.0	0.03452211590529575
id_2	1574913.0	-15732.615614206828

Took 4 sec. Last updated by anonymous at January 24 2017, 11:19:12 AM.

```
val matchData = parsed.as[MatchData] // as() : 새로운 DataSet를 리턴한다.
val scored = matchData.map(md => {
  (scoreMatchData(md), md.is_match)
}).toDF("score", "is_match")
crossTabs(scored, 4.0).show()
```

matchData: org.apache.spark.sql.Dataset[MatchData] = [id_1: int, id_2: int ... 10 more fields]

scored: org.apache.spark.sql.DataFrame = [score: double, is_match: boolean]

is_match	score
true	2087
false	6

Took 6 sec. Last updated by anonymous at January 24 2017, 11:19:15 AM.

Q & A